



Professional ML Project

Documentation: Student Dropout Risk Prediction

This document outlines the end-to-end plan for developing, deploying, and maintaining a machine learning model designed to predict student dropout risk. It follows a robust MLOps framework, integrating data versioning (DVC), experiment tracking (MLflow/W&B), and containerized deployment (Docker/FastAPI/AWS).

1. Project Repository Structure

The following structure ensures all artifacts—code, data, models, configurations, and documentation—are organized and version-controlled.

```
Student_dropout_risk/
├── .github/           # CI/CD workflows (GitHub Actions)
├── configs/           # YAML/JSON configs for pipeline
│   └── training_config.yaml  # Hyperparameters, feature lists, thresholds
├── data/
│   ├── raw/           # Original uncleaned data (DVC-tracked)
│   └── processed/      # Cleaned and engineered data (DVC-tracked)
├── models/           # Trained model artifacts (DVC/MLflow)
│   └── dropout_model.pkl
├── src/              # Source code for the entire project
│   ├── data/          # Scripts for data collection and initial loading
│   └── preprocessing/  # Cleaning and feature engineering scripts (e.g., encoders,
│                           scalers)
│       ├── models/     # Training, tuning, and evaluation code
│       ├── api/        # FastAPI application for serving predictions
│       │   └── app.py
│       ├── ui/         # Streamlit/Gradio UI code for a simple interface
│       └── monitoring/ # Scripts for model drift and latency checks
├── tests/            # Unit & integration tests
│   ├── test_preprocessing.py
│   ├── test_model.py
│   └── test_api.py
├── notebooks/        # Jupyter notebooks for EDA and experimentation
├── reports/          # Visual artifacts, evaluation reports, and model comparisons
├── experiments/      # MLflow/W&B logs and run metadata
├── .dvc/             # DVC configuration files
├── Dockerfile        # Docker container definition
└── requirements.txt   # Python dependencies
```

└─ README.md	# Project overview & instructions
└─ deployment.md	# Step-by-step deployment guide

2. End-to-End ML Project Roadmap & Tracking

This section details the nine phases of the project, including the specific actions taken for the Student Dropout Risk Prediction task and the professional tools used for tracking and reproducibility.

Phase 1: Problem Definition & Planning

Goal	Dropout Project Specifics	Success Tracking & Tools
Business Goal	Reduce student dropout rate by 20% through proactive intervention.	Define business goal, document in Confluence/Notion .
ML Type	Binary Classification (Dropout = 1, Continue = 0).	GitHub Wiki documentation.
Success Metric	F1-score or ROC-AUC (Target: ROC-AUC > 0.85).	Log criteria in configs/training_config.yaml .
Constraints	Predictions must be highly interpretable (explainable) for intervention teams.	Document assumptions & constraints in docs/constraints.md .

Phase 2: Data Collection & Versioning

Goal	Dropout Project Specifics	Success Tracking & Tools
Source	Open datasets (UCI Student Performance) or Synthetic generation.	Log source, date, and schema in data/schema.json .

Features	Demographics, Academics (Grades, Attendance), Behavior (Disciplinary records).	Ensure scripts are in <code>src/data/</code> .
Storage	Initial storage as CSV files.	Data stored in <code>data/raw/</code> .
Versioning	Use Data Version Control (DVC) to track the raw dataset.	Use DVC (<code>dvc add data/raw/</code>) for versioning.

Phase 3: Data Preprocessing & Cleaning

Goal	Dropout Project Specifics	Success Tracking & Tools
Cleaning	Handle missing values (median/mode imputation), duplicates, and outliers.	Keep scripts in <code>src/preprocessing/</code> .
Engineering	OneHotEncoding for categories; MinMaxScaler/StandardScaler for continuous features.	Use a reproducible pipeline (e.g., Scikit-Learn Pipeline).
Imbalance	Handle class imbalance using SMOTE on the training set or via model class weights .	Log transformations and class distribution using MLflow/W&B .
Splitting	Train/Test Split (80-20), ensuring stratification on the target variable.	Store processed data in <code>data/processed/</code> (DVC-tracked).
Testing	Validate preprocessing logic (e.g., no data leakage).	Write unit tests in <code>tests/test_preprocessing.py</code> .

Phase 4: Exploratory Data Analysis (EDA)

Goal	Dropout Project Specifics	Success Tracking & Tools
Visualization	Generate Correlation heatmap of features vs. dropout. Distribution plots for attendance and grades.	Save plots and summary in reports/eda_report.html .
Biases/Trends	Detect potential biases (e.g., income, gender). Analyze feature distributions.	Document findings in the EDA report.
Scripts	Keep analysis scripts versioned.	Analysis performed in notebooks/ and versioned via Git.

Phase 5: Model Selection & Training

Goal	Dropout Project Specifics	Success Tracking & Tools
Model	LightGBM Classifier (for speed and robust performance on structured data).	Training script in src/models/train_model.py .
Tuning	Hyperparameter search using GridSearchCV / Optuna for optimization.	MLflow / W&B to log experiments (hyperparameters, metrics, training time).

Code

```
import lightgbm as lgb
# ... (rest of the code)
model = lgb.LGBMClassifier(n_estimators=500, learning_rate=0.05, max_depth=7)
model.fit(X_train, y_train)
```

| Log training environment and dependencies. |

Phase 6: Model Evaluation

Goal	Dropout Project Specifics	Success Tracking & Tools
Metrics	Compare models based on ROC-AUC , F1-score, Precision, and Recall.	Log confusion matrix and ROC curves in MLflow / W&B .
Generalization	Use k-fold cross-validation during tuning to confirm model stability.	Save evaluation reports in <code>reports/model_comparison.md</code> .
Artifacts	Save the final, best-performing model.	Save model to <code>models/dropout_model.pkl</code> . Use DVC to version the model file.
Testing	Ensure the model output is correct given various inputs.	Write prediction tests in <code>tests/test_model.py</code> .

Phase 7: Deployment Preparation (API Development)

Goal	Dropout Project Specifics	Success Tracking & Tools
API Framework	FastAPI to create a high-performance, asynchronous prediction service.	API code in <code>src/api/app.py</code> .
Endpoint	Define a <code>POST /predict</code> endpoint that accepts a JSON payload of features and returns the dropout risk (0 or 1).	Document API specification in <code>docs/api_spec.yaml</code> .

API Code

```
from fastapi import FastAPI
import joblib
# ... (imports)
```

```
@app.post("/predict")
def predict(features: dict):
    # ... preprocessing and prediction logic
    prediction = model.predict(df)[0]
    return {"dropout_risk": int(prediction)}
```

Phase 8: Deployment & CI/CD

Phase 9: Monitoring & Maintenance

Goal	Dropout Project Specifics	Success Tracking & Tools
------	---------------------------	--------------------------

Monitoring	Track model drift, data drift, API latency, and error rates in production.	Use AWS CloudWatch for logs and Prometheus/Grafana for dashboarding. Monitoring logic in src/monitoring/ .
Data Logging	Log all prediction inputs and outputs (with timestamps) for future labeling and retraining.	Prediction logs saved to a dedicated S3 bucket.
Retraining Plan	Schedule quarterly retraining or trigger retraining when ROC-AUC drops by 5%.	Maintain schedule in docs/retraining_plan.md .
UI	Use Streamlit (code in src/ui/) for a simple, non-technical dashboard to visualize risk and monitoring metrics.	Log model updates and versions in the monitoring dashboard.

This comprehensive documentation covers the entire lifecycle of the ML project, from the initial business definition to production monitoring. Let me know if you would like to start drafting the code for a specific file, like [src/api/app.py](#) or [src/models/train_model.py](#)!